

UnsupervisedLearning - WEEK 3 - ReinforcementLearning - QUIZ 2 - ContinuousStateSpaces

Link: [UnsupervisedLearning - WEEK 3 - ReinforcementLearning - QUIZ 2 - ContinuousStateSpaces](#)

Your grade: 100%

Your latest: 100% • Your highest: 100% • To pass you need at least 80%. We keep your highest score.

Next item →

1 / 1 point

1. The Lunar Lander is a continuous state Markov Decision Process (MDP) because:

- ☒ The state contains numbers such as position and velocity that are continuous valued.
- ☐ The state-action value $Q(s, a)$ function outputs continuous valued numbers
- ☐ The state has multiple numbers rather than only a single number (such as position in the x -direction)
- ☐ The reward contains numbers that are continuous valued

✓ Correct
That's right!

1 / 1 point

2. In the learning algorithm described in the videos, we repeatedly create an artificial training set to which we apply supervised learning where the input $x = (s, a)$ and the target, constructed using Bellman's equations, is $y = \text{_____}$?

- ☐ $y = R(s)$
- ☐ $y = R(s')$ where s' is the state you get to after taking action a in state s
- ☒ $y = R(s) + \gamma \max_{a'} Q(s', a')$ where s' is the state you get to after taking action a in state s
- ☐ $y = \max_{a'} Q(s', a')$ where s' is the state you get to after taking action a in state s

✓ Correct

1 / 1 point

3. You have reached the final practice quiz of this class! What does that mean? (Please check all the answers, because all of them are correct!)

☒ What an accomplishment -- you made it!

☒ Correct

☒ You deserve to celebrate!

☒ Correct

☒ The DeepLearning.AI and Stanford Online teams would like to give you a round of applause!

☒ Correct

☒ Andrew sends his heartfelt congratulations to you!

☒ Correct